

Unit 3

Multiple Linear Regression

EL-GY 6143: INTRODUCTION TO MACHINE LEARNING

PROF. PEI LIU



NYU

TANDON SCHOOL
OF ENGINEERING

1



Learning Objectives

- ❑ Formulate a machine learning model as a multiple linear regression model.
 - Identify prediction vector and target for the problem.
- ❑ Write the regression model in matrix form. Write the feature matrix
- ❑ Compute the least-squares solution for the regression coefficients on training data.
- ❑ Derive the least-squares formula from minimization of the RSS

- ❑ Manipulate 2D arrays in python (indexing, stacking, computing shapes, ...)
- ❑ Compute the LS solution using python linear algebra and machine learning packages



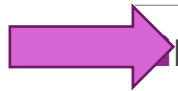
NYU

TANDON SCHOOL
OF ENGINEERING

2



Outline



Motivating Example: Understanding glucose levels in diabetes patients

- ❑ Multiple variable linear models
- ❑ Least squares solutions
- ❑ Computing the solutions in python
- ❑ Special case: Simple linear regression
- ❑ Extensions




NYU

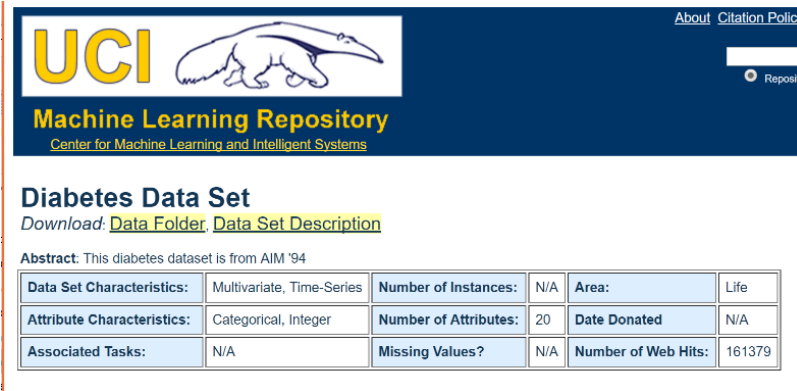
TANDON SCHOOL
OF ENGINEERING

Example: Blood Glucose Level

- ❑ Diabetes patients must monitor glucose level
- ❑ What causes blood glucose levels to rise and fall?
- ❑ Many factors
- ❑ We know mechanisms **qualitatively**
- ❑ But, **quantitative** models are difficult to obtain
 - Hard to derive from first principles
 - Difficult to model physiological process precisely
- ❑ Can machine learning help?



Data from AIM 94 Experiment



The screenshot shows the UCI Machine Learning Repository header with the UCI logo and a sloth illustration. Below the header, the title "Diabetes Data Set" is displayed, followed by download links for the "Data Folder" and "Data Set Description". An abstract states the dataset is from AIM '94. A table provides key statistics:

Data Set Characteristics:	Multivariate, Time-Series	Number of Instances:	N/A	Area:	Life
Attribute Characteristics:	Categorical, Integer	Number of Attributes:	20	Date Donated	N/A
Associated Tasks:	N/A	Missing Values?	N/A	Number of Web Hits:	161379

☐ Data collected as series of events

- Eating
- Exercise
- Insulin dosage

☐ Target variable glucose level monitored

Data Set Information:

Diabetes patient records were obtained from two sources: an electronic medical record system and a paper record. The electronic records are assigned to breakfast (08:00), lunch (12:00), dinner (18:00), and midnight (00:00). The paper records have more realistic time stamps.

Diabetes files consist of four fields per record. Each field is separated by a space.

File Names and format:

- (1) Date in MM-DD-YYYY format
- (2) Time in XX:YY format
- (3) Code
- (4) Value

The Code field is deciphered as follows:

- 33 = Regular insulin dose
- 34 = NPH insulin dose
- 35 = UltraLente insulin dose
- 48 = Unspecified blood glucose measurement
- 57 = Unspecified blood glucose measurement
- 58 = Pre-breakfast blood glucose measurement
- 59 = Post-breakfast blood glucose measurement
- 60 = Pre-lunch blood glucose measurement
- 61 = Post-lunch blood glucose measurement
- 62 = Pre-supper blood glucose measurement
- 63 = Post-supper blood glucose measurement



Demo on GitHub

□ All code is available in github:

https://github.com/pliugithub/MachineLearning/blob/master/unit03_mult_lin_reg/demo_glucose.ipynb

Demo: Predicting Glucose Levels using Multiple Linear Regression

In this demo, you will learn how to:

- Fit multiple linear regression models using python's sklearn package.
- Split data into training and test.
- Manipulating and visualizing multivariable arrays.

We first load the packages as usual.

```
import numpy as np
import matplotlib
import matplotlib.pyplot as plt
%matplotlib inline
```

Diabetes Data Example

To illustrate the concepts, we load the well-known diabetes data set. This dataset is included in the sklearn.data module and can be loaded as follows.

```
from sklearn import datasets, linear_model, preprocessing
```



NYU

TANDON SCHOOL
OF ENGINEERING

6



Using Google Colaboratory

❑ Two options for running the demos:


❑ Option 1:

- Clone the github repository to your local machine and run it using jupyter notebook
- Need to install all the software correctly

❑ Option 2: Run on the cloud in Google Colaboratory

❑ For Option 2:

- Go to <https://colab.research.google.com/>
- File->Open Notebook
- Select GitHub tab
- Enter github URL:
<https://github.com/sdrangan/introml>
- Select the unit03_mult_lin_reg/demo1_glucose.ipynb



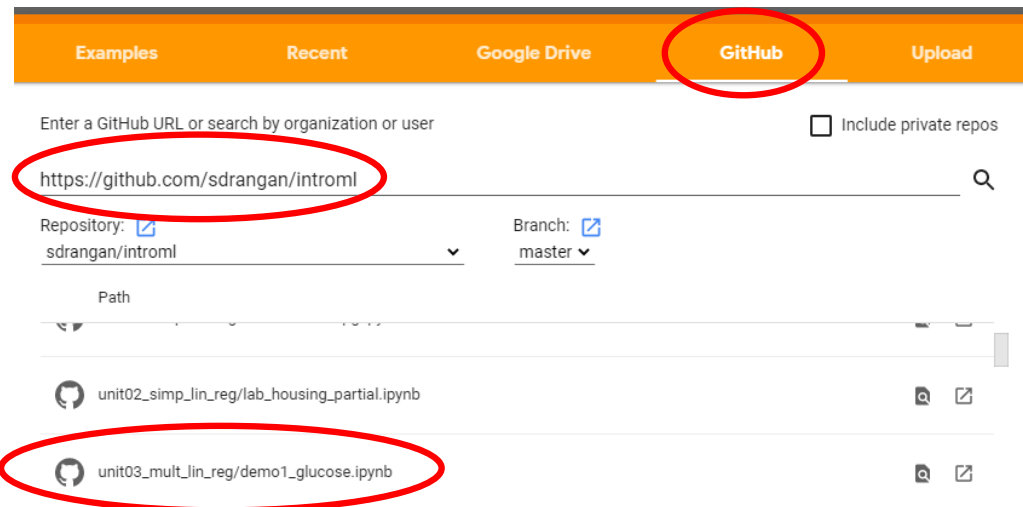
NYU

TANDON SCHOOL
OF ENGINEERING


7



Demo on Google Colab



Outline

- ❑ Motivating Example: Understanding glucose levels in diabetes patients
- ❑ Multiple variable linear models
- ❑ Least squares solutions
- ❑ Computing the solutions in python
-  ❑ Special case: Simple linear regression
- ❑ Extensions



NYU

TANDON SCHOOL
OF ENGINEERING

Simple vs. Multiple Regression

□ Simple linear regression: One predictor (feature)

- Scalar predictor x
- Linear model: $\hat{y} = \beta_0 + \beta_1 x$
- Can only account for one variable

□ Multiple linear regression: Multiple predictors (features)

- Vector predictor $\mathbf{x} = (x_1, \dots, x_k)$
- Linear model: $\hat{y} = \beta_0 + \beta_1 x_1 + \dots + \beta_k x_k$
- Can account for multiple predictors
- Turns into simple linear regression when $k = 1$



NYU

TANDON SCHOOL
OF ENGINEERING

10



Comparison to Single Variable Models

- We could compute models for each variable separately:

$$\begin{aligned}y &= a_1 + b_1 x_1 \\y &= a_2 + b_2 x_2 \\&\vdots\end{aligned}$$

- But, doesn't provide a way to account for joint effects
- Example: Consider three linear models to predicting longevity:
 - A: Longevity vs. some factor in diet (e.g. amount of fiber consumed)
 - B: Longevity vs. exercise
 - C: Longevity vs. diet AND exercise
 - What does C tell you that A and B do not?



Special Case: Single Variable

□ Suppose $k = 1$ predictor.

□ Feature matrix and coefficient vector:

$$A = \begin{bmatrix} 1 & x_1 \\ \vdots & \vdots \\ 1 & x_n \end{bmatrix}, \quad \beta = \begin{bmatrix} \beta_0 \\ \beta_1 \end{bmatrix}$$

□ LS soln: $\beta = \left(\frac{1}{N}A^T A\right)^{-1} \left(\frac{1}{N}A^T y\right) = P^{-1}r$

$$P = \begin{bmatrix} 1 & \bar{x} \\ \bar{x} & \bar{x}^2 \end{bmatrix}, \quad r = \begin{bmatrix} \bar{y} \\ \overline{xy} \end{bmatrix}$$

□ Obtain single variable solutions for coefficients (after some algebra):

$$\beta_1 = \frac{s_{xy}}{s_x^2}, \quad \beta_0 = \bar{y} - \beta_1 \bar{x}, \quad R^2 = \frac{s_{xy}^2}{s_x^2 s_y^2}$$



Loading the Data

```
: from sklearn import datasets, linear_model, preprocessing  
  
# Load the diabetes dataset  
diabetes = datasets.load_diabetes()  
X = diabetes.data  
y = diabetes.target
```

❑ scikit-learn package:

- Many methods for machine learning
- Datasets
- Will use throughout this class

❑ Diabetes dataset is one example

```
▶ print(diabetes.DESCR)
```

```
⦿ .. _diabetes_dataset:
```

Diabetes dataset

Ten baseline variables, age, sex, body mass index, average blood pressure, and six blood serum measurements were obtained for each of n = 442 diabetes patients, as well as the response of interest, a quantitative measure of disease progression one year after baseline.

****Data Set Characteristics:****

:Number of Instances: 442

:Number of Attributes: First 10 columns are numeric predictive values

:Target: Column 11 is a quantitative measure of disease progression one year after baseline

:Attribute Information:

- Age
- Sex
- Body mass index
- Average blood pressure
- S1
- S2
- S3
- S4
- S5
- S6



NYU

TANDON SCHOOL
OF ENGINEERING

Simple Linear Regression for Diabetes Data

```
ym = np.mean(y)
syy = np.mean((y-ym)**2)
Rsqr = np.zeros(natt)
for k in range(natt):
    xm = np.mean(X[:,k])
    sxy = np.mean((X[:,k]-xm)*(y-ym))
    sxx = np.mean((X[:,k]-xm)**2)
    Rsqr[k] = (sxy)**2/sxx/syy

print("{0:2d}  Rsqr={1:f}".format(k, Rsqr[k]))
```

```
0  Rsqr=0.035302
1  Rsqr=0.001854
2  Rsqr=0.343924
3  Rsqr=0.194908
4  Rsqr=0.044954
5  Rsqr=0.030295
6  Rsqr=0.155859
7  Rsqr=0.185290
8  Rsqr=0.320224
9  Rsqr=0.146294
```

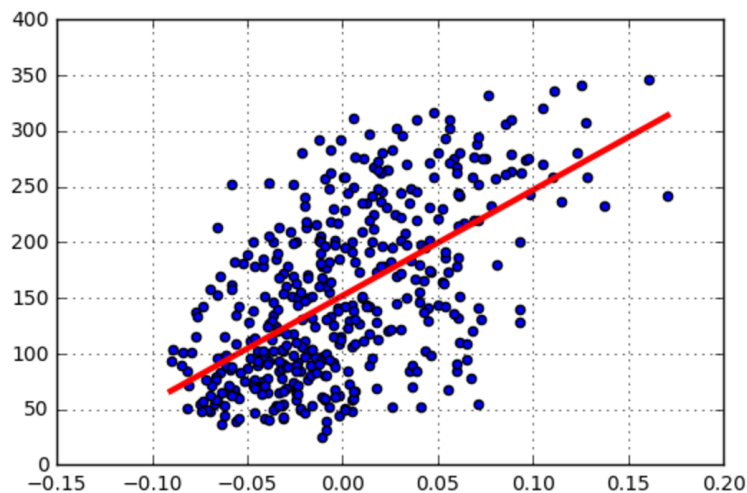
Best individual variable

- ❑ Try a fit of each variable individually
- ❑ Compute R_k^2 coefficient for each variable
- ❑ Use formula on previous slide
- ❑ “Best” individual variable is a poor fit
 - $R_k^2 \approx 0.34$



Scatter Plot

- ❑ No one variable explains glucose well
- ❑ Multiple linear regression could be much better



```
# Find the index of the single variable with the best R^2
imax = np.argmax(Rsq)


# Regression Line over the range of x values
xmin = np.min(X[:,imax])
xmax = np.max(X[:,imax])
ymin = beta0[imax] + beta1[imax]*xmin
ymax = beta0[imax] + beta1[imax]*xmax
plt.plot([xmin,xmax], [ymin,ymax], 'r-', linewidth=3)

# Scatter plot of points
plt.scatter(X[:,imax],y)
plt.grid()
```



Outline

- ❑ Motivating Example: Understanding glucose levels in diabetes patients

- ❑ Multiple variable linear models

- ❑ Least squares solutions

- ❑ Computing the solutions in python

- ❑ Special case: Simple linear regression

- ❑ Extensions



NYU

TANDON SCHOOL
OF ENGINEERING

16



Finding a Mathematical Model

Attributes

x_1 : Age
 x_2 : Sex
 x_3 : BMI
 x_4 : BP
 x_5 : S1
⋮
 x_{10} : S6



Target

y = Glucose level

$$y \approx \hat{y} = f(x_1, \dots, x_{10})$$

□ **Goal:** Find a function to predict glucose level from the 10 attributes

□ **Problem:** Several attributes

- Need a **multi-variable function**



NYU

TANDON SCHOOL
OF ENGINEERING

Matrix Representation of Data

□ Data is a **matrix** and a target **vector**

□ n samples:

- One sample per row

□ k features / attributes / predictors:

- One feature per column

$$X = \begin{bmatrix} x_{11} & \cdots & x_{1k} \\ \vdots & \ddots & \vdots \\ x_{n1} & \cdots & x_{nk} \end{bmatrix}$$

Attributes

Target vector

$$y = \begin{bmatrix} y_1 \\ \vdots \\ y_n \end{bmatrix}$$

Samples

□ This example:

- y_i = blood glucose measurement of i -th sample
- $x_{i,j}$: j -th feature of i -th sample
- $\mathbf{x}_i^T = [x_{i,1}, x_{i,2}, \dots, x_{i,k}]$: feature or predictor vector
- i -th sample contains \mathbf{x}_i, y_i



In class exercise

In class exercise: What are the (normalized) ages of the first 5 subjects?

```
[18] # TODO
```

In-class exercise: Print the attributes S1-S3 for subjects 10-15

```
[ ] # TODO
```

Double-click (or enter) to edit

In class exercise: Create a scatter plot of the target variable, y vs. the BMI. Does there seem to be a relation? What about y vs. the age? Which is a better predictor?

```
[19] # TODO
```



NYU


TANDON SCHOOL
OF ENGINEERING

19



Outline

- ❑ Motivating Example: Understanding glucose levels in diabetes patients

-  Multiple variable linear models

- ❑ Least squares solutions

- ❑ Computing the solutions in python

- ❑ Special case: Simple linear regression

- ❑ Extensions



NYU

TANDON SCHOOL
OF ENGINEERING

20



Multivariable Linear Model for Glucose

Attributes

Age, Sex, BMI, BP, S1, ..., S6
 $\mathbf{x} = [x_1, \dots, x_{10}]$

$$y \approx \hat{y} = f(x_1, \dots, x_{10})$$



Target

y = Glucose level

□ **Goal:** Find a function to predict glucose level from the 10 attributes

□ **Linear Model:** Assume glucose is a **linear function** of the predictors:

$$[\text{glucose}] \approx [\text{prediction}] = \beta_0 + \beta_1[\text{Age}] + \dots + \beta_4[\text{BP}] + \beta_5[\text{S1}] + \dots + \beta_{10}[\text{S6}]$$

□ **General form:**

$$\begin{array}{ccccccc} y & \approx & \hat{y} & = & \beta_0 + & \underbrace{\beta_1 x_1 + \dots + \beta_4 x_4 + \beta_5 x_5 + \dots + \beta_{10} x_{10}}_{\text{10 Features}} \\ \text{Target} & & & \uparrow & & & \\ & & & \text{Intercept} & & & \end{array}$$



Multiple Variable Linear Model

❑ Vector of **features**: $\mathbf{x} = [x_1, \dots, x_k]$

- k features (also known as predictors, independent variable, attributes, covariates, ...)

❑ Single **target variable** y

- What we want to predict

❑ Linear model: Make a **prediction** \hat{y}

$$y \approx \hat{y} = \beta_0 + \beta_1 x_1 + \dots + \beta_k x_k$$

❑ **Data** for training

- Samples are (\mathbf{x}_i, y_i) , $i=1,2,\dots,n$.
- Each sample has a vector of features: $\mathbf{x}_i = [x_{i1}, \dots, x_{ik}]$ and scalar target y_i

❑ **Problem**: Learn the best **coefficients** $\boldsymbol{\beta} = [\beta_0, \beta_1, \dots, \beta_k]$ from the training data



Example: Heart Rate Increase

□ **Linear Model:** $[\text{HR increase}] \approx \beta_0 + \beta_1[\text{mins exercise}] + \beta_2[\text{exercise intensity}]$

□ **Data:**

Subject number	HR before	HR after	Mins on treadmill	Speed (min/km)	Days exercise / week
123	60	90	1	5.2	3
456	80	110	2	4.1	1
789	70	130	5	3.5	2
⋮	⋮	⋮	⋮	⋮	⋮



Measuring fitness of athletes

<https://www.mercurynews.com/2017/10/29/4851089/>



NYU

TANDON SCHOOL
OF ENGINEERING

23



Why Use a Linear Model?

❑ Many natural phenomena have linear relationship

❑ Predictor has small variation

- Suppose $y = f(x)$
- If variation of x is small around some value x_0 , then

$$y \approx f(x_0) + f'(x_0)(x - x_0) = \beta_0 + \beta_1 x,$$

$$\beta_0 = f(x_0) - f'(x_0)x_0, \quad \beta_1 = f'(x_0)$$

❑ Simple to compute

❑ Easy to interpret relation

- Coefficient β_j indicates the importance of feature j for the target.

❑ Advanced: Gaussian random variables:

- If two variables are jointly Gaussian, the optimal predictor of one from the other is linear predictor



NYU

TANDON SCHOOL
OF ENGINEERING

24



Matrix Review

□ Consider

$$A = \begin{bmatrix} 1 & 2 \\ 3 & 4 \\ 5 & 6 \end{bmatrix}, \quad B = \begin{bmatrix} 2 & 0 \\ 3 & 2 \end{bmatrix}, \quad x = \begin{bmatrix} 2 \\ 3 \end{bmatrix},$$

□ Compute (computations on the board):

- Matrix vector multiply: Ax
- Transpose: A^T
- Matrix multiply: AB
- Solution to linear equations: Solve for u : $x = Bu$
- Matrix inverse: B^{-1}



Slopes, Intercept and Inner Products

□ Model with **coefficients** β : $\hat{y} = \beta_0 + \beta_1 x_1 + \cdots + \beta_k x_k$

□ Sometimes use **weight bias version**:

$$\hat{y} = b + w_1 x_1 + \cdots + w_k x_k$$

- $b = \beta_0$: **Bias** or **intercept**
- $\mathbf{w} = \beta_{1:k} = [\beta_1, \dots, \beta_k]$: **Weights** or **slope vector**

□ Can write either with **inner product**:

$$\hat{y} = \beta_0 + \beta_{1:k} \cdot \mathbf{x}$$

or

$$\hat{y} = b + \mathbf{w} \cdot \mathbf{x}$$

□ Inner product:

- $\mathbf{w} \cdot \mathbf{x} = \sum_{j=1}^k w_j x_j$
- Will use alternate notation: $\mathbf{w}^T \mathbf{x} = \langle \mathbf{w}, \mathbf{x} \rangle$



Matrix Form of Linear Regression

□ Data: $(x_i, y_i), i = 1, \dots, n$

□ Predicted value for i -th sample: $\hat{y}_i = \beta_0 + \beta_1 x_{i1} + \dots + \beta_k x_{ik}$

□ Matrix form

$$\hat{\mathbf{y}} \text{ a } n \text{ predicted values } \left\{ \begin{bmatrix} \hat{y}_1 \\ \hat{y}_2 \\ \vdots \\ \hat{y}_n \end{bmatrix} \right\} = \underbrace{\begin{bmatrix} 1 & x_{11} & \cdots & x_{1k} \\ 1 & x_{21} & \cdots & x_{2k} \\ \vdots & \vdots & \cdots & \vdots \\ 1 & x_{n1} & \cdots & x_{nk} \end{bmatrix}}_{\mathbf{A} \text{ a } n \times k \text{ feature matrix}} \underbrace{\begin{bmatrix} \beta_0 \\ \beta_1 \\ \vdots \\ \beta_k \end{bmatrix}}_{\boldsymbol{\beta} \text{ with } p = k + 1 \text{ coefficient vector}}$$

□ Matrix equation:

$$\hat{\mathbf{y}} = \mathbf{A} \boldsymbol{\beta}$$



In-Class Exercise

Consider a linear model:

$$[\text{HR increase}] \approx \beta_0 + \beta_1[\text{mins exercise}] + \beta_2[\text{exercise intensity}].$$

We are given the following data: Only the first three rows and the final entry are shown.


Subject number	HR before	HR after	Mins on treadmill	Speed (min/km)	Days exercise / week
123	60	90	1	5.2	3
456	80	110	2	4.1	1
789	70	130	5	3.5	2
⋮	⋮	⋮	⋮	⋮	⋮
283	75	100	1	4.8	0

100 subjects

- Q1: What is the feature matrix A and target vector y . What are their dimensions?
 - Fill in only the values from the first three rows and the last row
- Q2. Suppose that after training, we find parameters $\beta = [0, 15, 3]$. If the initial HR is 70 bpm, what is the predicted HR after 2 minutes of exercise at 5 km/hr.



Outline

- ❑ Motivating Example: Understanding glucose levels in diabetes patients
- ❑ Multiple variable linear models
- ❑ Least squares solutions
- ❑ Computing the solutions in python
- ❑ Special case: Simple linear regression
- ❑ Extensions



NYU

TANDON SCHOOL
OF ENGINEERING

29



Least Squares Model Fitting

□ How do we select parameters $\boldsymbol{\beta} = (\beta_0, \dots, \beta_k)$?

□ Define $\hat{y}_i = \beta_0 + \beta_1 x_{i1} + \dots + \beta_k x_{ik}$

- Predicted value on sample i for parameters $\boldsymbol{\beta} = (\beta_0, \dots, \beta_k)$

□ Define average **residual sum of squares**:

$$\text{RSS}(\boldsymbol{\beta}) := \sum_{i=1}^n (y_i - \hat{y}_i)^2$$

- Note that \hat{y}_i is implicitly a function of $\boldsymbol{\beta} = (\beta_0, \dots, \beta_k)$
- Also called the sum of **squared residuals** (SSR) and **sum of squared errors** (SSE)

□ **Least squares solution**: Find $\boldsymbol{\beta}$ to minimize RSS.



Variants of RSS

❑ Often use some variant of RSS

◦ Note: these are not standard

❑ Residual sum of squares: $RSS = \sum_{i=1}^n (y_i - \hat{y}_i)^2$

❑ RSS per sample or Mean Squared Error:

$$MSE = \frac{RSS}{n} = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2$$

❑ Normalized RSS or Normalized MSE:

$$\frac{RSS/n}{s_y^2} = \frac{MSE}{s_y^2} = \frac{\sum_{i=1}^n (y_i - \hat{y}_i)^2}{\sum_{i=1}^n (y_i - \bar{y})^2}$$



Finding Parameters via Optimization

A general ML recipe

General ML problem

- ❑ Pick a **model** with **parameters**
- ❑ Get **data**
- ❑ Pick a **loss function**
 - Measures goodness of fit model to data
 - Function of the parameters

Multiple linear regression

- ➡ Linear model: $\hat{y} = \beta_0 + \beta_1 x_1 + \dots + \beta_k x_k$
- ➡ Data: $(x_i, y_i), i = 1, 2, \dots, n$
- ➡ Loss function:
$$RSS(\beta_0, \dots, \beta_k) := \sum (y_i - \hat{y}_i)^2$$
- ➡ Find parameters that **minimizes** loss ➡ Select $\boldsymbol{\beta} = (\beta_0, \dots, \beta_k)$ to minimize $RSS(\boldsymbol{\beta})$



RSS as a Vector Norm

□ RSS is given by sum:


$$\text{RSS} = \sum_{i=1}^n (y_i - \hat{y}_i)^2$$

□ Define **norm** of a vector:

- $\|x\| = (x_1^2 + \dots + x_r^2)^{1/2}$
- Standard Euclidean norm.
- Sometimes called ℓ -2 norm. ℓ is for Lebesgue

□ Write RSS in vector form:

$$\text{RSS} = \|y - \hat{y}\|^2$$



NYU

TANDON SCHOOL
OF ENGINEERING

Least Squares Solution

□ Consider cost function of the RSS:

$$\text{RSS}(\boldsymbol{\beta}) = \sum_{i=1}^n (y_i - \hat{y}_i)^2, \quad \hat{y}_i = \sum_{j=0}^p A_{ij}\beta_j$$

- Vector $\boldsymbol{\beta}$ that minimizes RSS called the **least-squares** solution

□ **Least squares solution**: The vector $\boldsymbol{\beta}$ that minimizes the RSS is:

$$\hat{\boldsymbol{\beta}} = (\mathbf{A}^T \mathbf{A})^{-1} \mathbf{A}^T \mathbf{y}$$

- Can compute the best coefficient vector analytically
- Just solve a linear set of equations
- Will show the proof below



Proving the LS Formula

□ **Least squares formula:** The vector $\boldsymbol{\beta}$ that minimizes the RSS is:

$$\hat{\boldsymbol{\beta}} = (\mathbf{A}^T \mathbf{A})^{-1} \mathbf{A}^T \mathbf{y}$$

□ To prove this formula, we will:

- Review gradients of multi-variable functions
- Compute gradient $\nabla RSS(\boldsymbol{\beta})$
- Solve $\nabla RSS(\boldsymbol{\beta}) = 0$



NYU

TANDON SCHOOL
OF ENGINEERING

35



Gradients of Multi-Variable Functions

Consider scalar valued function of a vector: $f(\boldsymbol{\beta}) = f(\beta_1, \dots, \beta_n)$

Gradient is the column vector:

$$\nabla f(\boldsymbol{\beta}) = \begin{bmatrix} \partial f(\boldsymbol{\beta}) / \partial \beta_1 \\ \vdots \\ \partial f(\boldsymbol{\beta}) / \partial \beta_n \end{bmatrix}$$



Represents direction of maximum increase

At a local minima or maxima: $\nabla f(\boldsymbol{\beta}) = 0$

- Solve n equations and n unknowns

Ex: $f(\beta_1, \beta_2) = \beta_1 \sin \beta_2 + \beta_1^2 \beta_2$.

- Compute $\nabla f(\boldsymbol{\beta})$. Solution on board



NYU

TANDON SCHOOL
OF ENGINEERING

36



Proof of the LS Formula

□ Consider cost function of the RSS:

$$RSS = \sum_{i=1}^n (y_i - \hat{y}_i)^2, \quad \hat{y}_i = \sum_{j=0}^p A_{ij} \beta_j$$

◦ Vector $\boldsymbol{\beta}$ that minimizes RSS called the **least-squares** solution

□ Compute partial derivatives via chain rule: $\frac{\partial RSS}{\partial \beta_j} = -2 \sum_{i=1}^n (y_i - \hat{y}_i) A_{ij}, j = 1, 2, \dots, k$

□ Matrix form: $RSS = \|A\boldsymbol{\beta} - \mathbf{y}\|^2, \nabla RSS = -2A^T(\mathbf{y} - A\boldsymbol{\beta})$

□ Solution: $A^T(\mathbf{y} - A\boldsymbol{\beta}) = 0 \rightarrow \boldsymbol{\beta} = (A^T A)^{-1} A^T \mathbf{y}$ (least squares solution of equation $A\boldsymbol{\beta} = \mathbf{y}$)

□ Minimum RSS: $RSS = \mathbf{y}^T [I - A(A^T A)^{-1} A^T] \mathbf{y}$

◦ Proof on the board



NYU

TANDON SCHOOL
OF ENGINEERING

37



LS Solution via Auto-Correlation Functions

- Each data sample has a linear feature vector:

$$A_i = (A_{i0}, \dots, A_{ik}) = (1, x_{i1}, \dots, x_{ik})$$

- Define sample **auto-correlation** matrix and **cross-correlation** vector:

- $R_{AA} = \frac{1}{n} A^T A$, $R_{AA}(\ell, m) = \frac{1}{n} \sum_{i=1}^n A_{i\ell} A_{im}$ (correlation of feature ℓ and feature m)
- $R_{Ay} = \frac{1}{n} A^T y$, $R_{yA}(\ell) = \frac{1}{n} \sum_{i=1}^n A_{i\ell} y_i$ (correlation of feature ℓ and target)

- Least squares solution is: $\beta = R_{AA}^{-1} R_{Ay}$



Mean Removed Form of the LS Solution

□ Often useful to remove mean from data before fitting

□ Sample mean: $\bar{y} = \frac{1}{N} \sum_{i=1}^N y_i$, $\bar{x}_j = \frac{1}{N} \sum_{i=1}^N x_{ij}$, $\bar{\mathbf{x}} = [\bar{x}_1, \dots, \bar{x}_k]$

□ Defined **mean removed data**: $\tilde{X}_{ij} = x_{ij} - \bar{x}_j$, $\tilde{y}_i = y_i - \bar{y}$

□ Sample **covariance** matrix and **cross-covariance** vector:

$$\circ S_{xx}(\ell, m) = \frac{1}{N} \sum_{i=1}^N (x_{i\ell} - \bar{x}_\ell)(x_{im} - \bar{x}_m), \quad S_{xx} = \frac{1}{N} \tilde{\mathbf{X}}^T \tilde{\mathbf{X}}$$

$$\circ S_{xy}(\ell) = \frac{1}{N} \sum_{i=1}^N (x_{i\ell} - \bar{x}_\ell)(y_i - \bar{y}), \quad S_{xy} = \frac{1}{N} \tilde{\mathbf{X}}^T \tilde{\mathbf{y}}$$

□ **Mean-Removed form** of the least squares solution:

$$\hat{y} = \boldsymbol{\beta}_{1:k} \cdot \mathbf{x} + \beta_0, \quad \boldsymbol{\beta}_{1:k} = S_{xx}^{-1} S_{xy}, \quad \beta_0 = \bar{y} - \boldsymbol{\beta}_{1:k} \cdot \bar{\mathbf{x}}$$

□ Proof: On board



R^2 : Goodness of Fit

□ Multiple variable **coefficient of determination**:

$$R^2 = \frac{s_y^2 - MSE}{s_y^2} = 1 - \frac{MSE}{s_y^2}$$

- $MSE = \frac{RSS}{n} = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2$
- Sample variance is: $s_y^2 = \frac{1}{n} \sum_{i=1}^n (y_i - \bar{y})^2$, $\bar{y} = \frac{1}{n} \sum_{i=1}^n y_i$,

□ **Interpretation**:

- $\frac{MSE}{s_y^2} = \frac{\text{Error with linear predictor}}{\text{Error predicting by mean}}$
- R^2 = fraction of variance reduced or “explained” by the model.

□ On the training data (not necessarily on the test data):

- $R^2 \in [0,1]$ always
- $R^2 \approx 1 \Rightarrow$ linear model provides a good fit
- $R^2 \approx 0 \Rightarrow$ linear model provides a poor fit



In-Class Exercise


We are given the following data

Sample number	Target y_i	Feature 1 x_{i1}	Feature 2 x_{i2}
1	3.0	0	1
2	5.0	2	3
3	9.0	4	8
4	10.0	6	10

- Q1. Write the equations to solve for the linear model using all four data points
 - Write the feature matrix and the equations for coefficients.
 - Do not solve them (you would need a computer)
- Q2. Can you find parameters that exactly fits the first three data points?
 - Just state if such parameters exist. You do not need to find them.



Outline

- ❑ Motivating Example: Understanding glucose levels in diabetes patients
- ❑ Multiple variable linear models
- ❑ Least squares solutions
- ❑ Computing the solutions in python
- ❑ Special case: Simple linear regression
- ❑ Extensions



NYU

TANDON SCHOOL
OF ENGINEERING

42



Arrays and Vector in Python and MATLAB

□ There are some key differences between MATLAB and Python that you need to get used to

□ MATLAB

- All arrays are at least 2 dimensions
- Vectors are $1 \times N$ (row vectors) or $N \times 1$ (column) vectors
- Matrix vector multiplication syntax depends if vector is on left or right: $x' * A$ or $A * x$

□ Python:

- Arrays can have 1, 2, 3, ... dimension
- Vectors can be 1D arrays; matrices are generally 2D arrays
- Vectors that are 1D arrays are neither row nor column vectors
- If x is 1D and A is 2D, then left and right multiplication are the same: $x \cdot \text{dot}(A)$ and $A \cdot \text{dot}(x)$

□ **Lecture notes:** We will generally treat x and x^T the same.

- Can write $x = (x_1, \dots, x_N)$ and still multiply by a matrix on left or right



Fitting Using sklearn

```
: ns_train = 300
  ns_test = nsamp - ns_train
  X_tr = X[:ns_train,:]
  y_tr = y[:ns_train]
```

- ❑ Return to diabetes data example
- ❑ All code in demo
- ❑ Divide data into two portions:
 - Training data: First 300 samples
 - Test data: Remaining 142 samples
- ❑ Train model on training data.
- ❑ Test model (i.e. measure RSS) on test data
- ❑ Reason for splitting data discussed next lecture.



Manually Computing the Solution

```
ones = np.ones((ns_train,1))  
A = np.hstack((ones,X_tr))
```

```
out = np.linalg.lstsq(A,y_tr)  
beta = out[0]
```

❑ Use numpy linear algebra routine to solve
$$\beta = (A^T A)^{-1} A^T y$$

❑ Common mistake:

- Compute matrix inverse $P = (A^T A)^{-1}$,
- Then compute $\beta = P A^T y$
- Full matrix inverse is VERY slow. Not needed.
- Can directly solve linear system: $A \beta = y$
- Numpy has routines to solve this directly



NYU

TANDON SCHOOL
OF ENGINEERING

45



Calling the sklearn Linear Regression method

```
regr = linear_model.LinearRegression()  
regr.fit(X_tr,y_tr)
```


```
X_test = X[ns_train:,:]  
y_test = y[ns_train:]  
y_test_pred = regr.predict(X_test)  
RSS_test = np.mean((y_test_pred-y_test)**2)/(np.std(y_test)**2)  
Rsqr_test = 1-RSS_test  
print("RSS per sample = {0:f}".format(RSS_test))  
print("R^2 = {0:f}".format(Rsqr_test))
```

```
RSS per sample = 0.492801  
R^2 = 0.507199
```

We see that the model predicts new samples almost as well as it did the training samples

```
plt.scatter(y_test,y_test_pred)  
plt.plot([0,350],[0,350], 'r')  
plt.xlabel('Actual')  
plt.ylabel('Predicted')  
plt.grid()
```

- ❑ Construct a linear regression object
- ❑ Run it on the training data
- ❑ Predict values on the test data



In-Class Exercise

In-Class Simple Exercise


You are given target values `y` and features `x1` and `x2` below. Fit the model on the first 4 data points and test the model on the fifth data point. You may want to use the following steps

- Construct the training data `X_tr, y_tr`
- Create a regression object `regr = linear_model.LinearRegression()`
- Fit the model with the `regr.fit()` method
- Predict the value on the test value with the `regr.predict()`

```
: ▶ 1 x1 = np.array([0,1,3,5,4])
    2 x2 = np.array([0,0.7, 4.3, 15.1, 13.2])
    3 y = np.array([-2, -0.9, 1.5, 18, 13])
    4
    5 # TODO
```



Outline

- ❑ Motivating Example: Understanding glucose levels in diabetes patients
- ❑ Multiple variable linear models
- ❑ Least squares solutions
- ❑ Computing in python
- ❑ Extensions



NYU

TANDON SCHOOL
OF ENGINEERING

48



Transformed Linear Models

□ Standard linear model: $\hat{y} = \beta_0 + \beta_1 x_1 + \dots + \beta_d x_d$

□ Linear model may be too restrictive

- Relation between x and y can be **nonlinear**

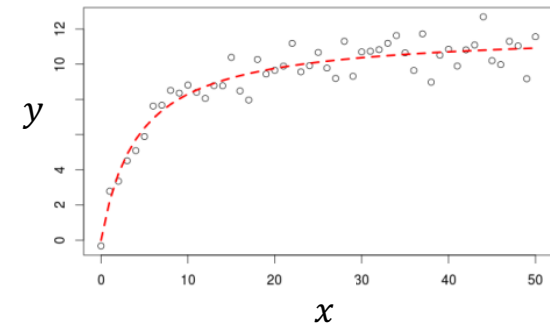
□ Useful to look at models in **transformed form**:

$$\hat{y} = \beta_1 \phi_1(\mathbf{x}) + \dots + \beta_p \phi_p(\mathbf{x})$$

- Each function $\phi_j(\mathbf{x}) = \phi_j(x_1, \dots, x_d)$ is called a **basis function**
- Each basis function may be nonlinear and a function of multiple variables

□ Can write in vector form: $\hat{y} = \boldsymbol{\phi}(\mathbf{x}) \cdot \boldsymbol{\beta}$

- $\boldsymbol{\phi}(\mathbf{x}) = [\phi_1(\mathbf{x}), \dots, \phi_p(\mathbf{x})]$, $\boldsymbol{\beta} = [\beta_1, \dots, \beta_p]$



Fitting Transformed Linear Models

- Consider transformed linear model

$$\hat{y} = \beta_1 \phi_1(x) + \cdots + \beta_p \phi_p(x)$$

- We can fit this model exactly as before

- Given data $(x_i, y_i), i = 1, \dots, N$
- Want to fit the model from the transformed variables $\phi_j(x)$ to target y
- Define the transformed matrix:

$$A = \begin{bmatrix} \phi_1(x_1) & \cdots & \phi_p(x_1) \\ \vdots & \vdots & \vdots \\ \phi_1(x_N) & \cdots & \phi_p(x_N) \end{bmatrix}$$

- Predictions: $\hat{y} = A\beta$
- Least squares fit $\hat{\beta} = (A^T A)^{-1} A^T y$

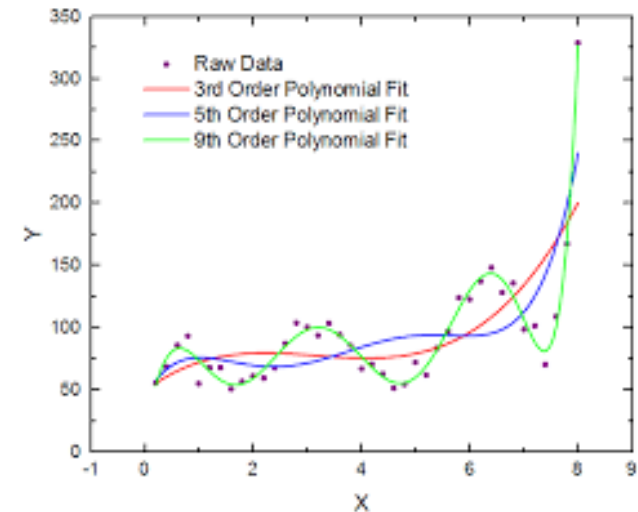


Example: Polynomial Fitting

- Suppose y only depends on a single variable x ,
- Want to fit a polynomial model
 - $y \approx \beta_0 + \beta_1 x + \dots \beta_d x^d$
- Given data $(x_i, y_i), i = 1, \dots, n$
- Take basis functions $\phi_j(x) = x^j, j = 0, \dots, d$
- Transformed model: $\hat{y} = \beta_0 \phi_0(x) + \dots + \beta_d \phi_d(x)$
- Transformed matrix is:

$$A = \begin{bmatrix} 1 & x_1 & \dots & x_1^d \\ \vdots & \vdots & \dots & \vdots \\ 1 & x_n & \dots & x_n^d \end{bmatrix}, \quad \beta = \begin{bmatrix} \beta_0 \\ \vdots \\ \beta_d \end{bmatrix}$$

- $p = d + 1$ transformed features from 1 original feature
- Will discuss how to select d in the next lecture



Other Nonlinear Examples

□ **Multinomial model:** $\hat{y} = a + b_1x_1 + b_2x_2 + c_1x_1^2 + c_2x_1x_2 + c_3x_2^2$

- Contains all second order terms
- Define parameter vector $\beta = [a, b_1, b_2, c_1, c_2, c_3]$
- Transformed vector $\phi(x_1, x_2) = [1, x_1, x_2, x_1^2, x_1x_2, x_2^2]$
- Note that the features are nonlinear functions of $\mathbf{x} = [x_1, x_2]$

□ **Exponential model:** $\hat{y} = a_1e^{-b_1x} + \dots + a_de^{-b_dx}$

- If the parameters b_1, \dots, b_d are fixed, then the model is linear in the parameters a_1, \dots, a_d
- Parameter vector $\beta = [a_1, \dots, a_d]$
- Transformed vector $\phi(x) = [e^{-b_1x}, \dots, e^{-b_dx}]$
- But, if the parameters b_1, \dots, b_d are not fixed, the model is nonlinear in b_1, \dots, b_d



Linear Models via Re-Parametrization

□ Sometimes models can be made into a linear model via re-parametrization

□ Example: Consider the model: $\hat{y} = Ax_1(1 + Be^{-x_2})$

◦ Parameters (A, B)

□ This is **nonlinear** in (A, B) due to the product AB : $\hat{y} = Ax_1 + ABx_1e^{-x_2}$

□ But, we can define a new set of parameters:

◦ $\beta_1 = A$ and $\beta_2 = AB$

□ Then, $\hat{y} = \beta_1x_1 + \beta_2x_1e^{-x_2}$

□ Basis functions: $\phi(x_1, x_2) = [x_1, x_1e^{-x_2}]$

□ After we solve for β_1, β_2 we can recover A, B via inverting the equations:

$$A = \beta_1, \quad B = \frac{\beta_2}{A}$$



Example: Learning Linear Systems

□ Linear system: $y_k = a_1 y_{k-1} + \dots + a_m y_{k-m} + b_0 x_k + \dots + b_n x_{k-n} + w_k$

□ Transfer function: $H(z) = \frac{b_0 + \dots + b_n z^{-n}}{1 - a_1 z^{-1} - \dots - a_m z^{-m}}$

□ Given input sequence and output sequence for T samples,

How do we determine $\beta = (a_1, \dots, a_m, b_0, \dots, b_n)^T$.

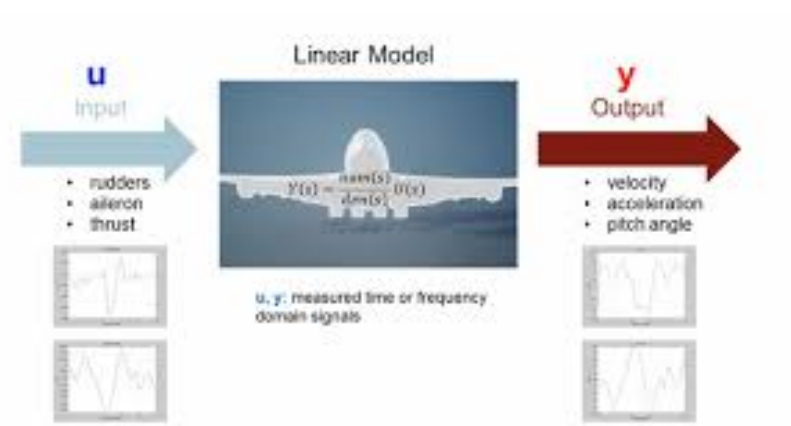
□ Can be solved using linear regression!

□ Write $y = A\beta + w$ and define A, y

- See homework problem

□ Many applications

- Learning dynamics in robots / mechanical systems
- Modeling responses in neural systems
- Stock market time series
- Speech modeling. Fit a model each 25 ms.



NYU

TANDON SCHOOL
OF ENGINEERING

54



One Hot Encoding

□ Suppose that one feature x_j is a **categorical** variable

□ Ex: Predict the price of a car, y , given model x_1 and interior space x_2

- Suppose there are 3 different models of a car (Ford, BMW, GM)
- Bad idea: Arbitrarily assign an index to each possible car model
- Can give unreasonable relations

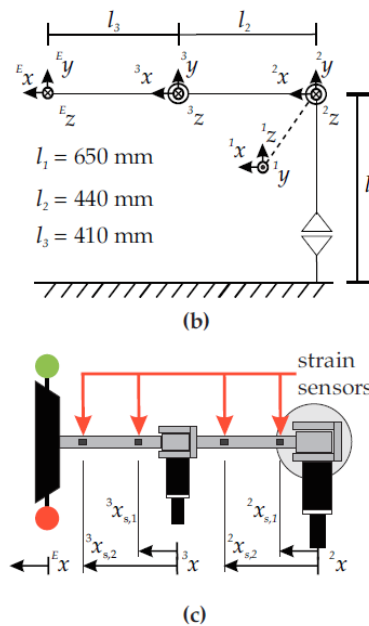
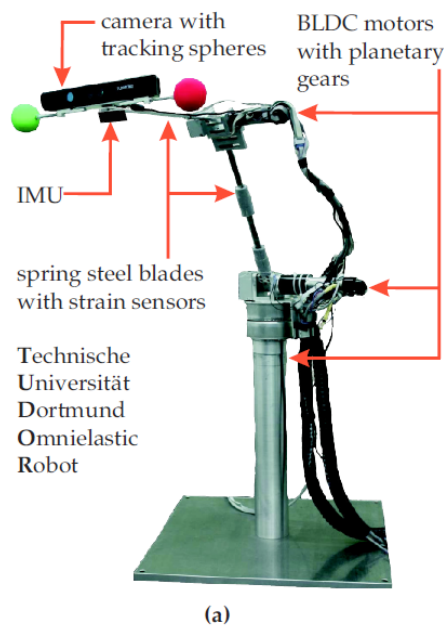
Model	ϕ_1	ϕ_2	ϕ_3
Ford	1	0	0
BMW	0	1	0
GM	0	0	1

□ One-hot encoding example:

- With 3 possible categories, represent x_1 using 3 binary features (ϕ_1, ϕ_2, ϕ_3)
- Model: $y = \beta_0 + \beta_1\phi_1 + \beta_2\phi_2 + \beta_3\phi_3 + \beta_4x_2$
- Essentially obtain 3 different models:
 - Ford: $y = \beta_0 + \beta_1 + \beta_4x_2$
 - BMW: $y = \beta_0 + \beta_2 + \beta_4x_2$
 - GM: $y = \beta_0 + \beta_3 + \beta_4x_2$
- Allows different intercepts (or mean values) for different categories!



Lab: Robot Calibration



□ Predict the current draw

- Needed to predict power consumption

□ Predictors:

- Joint angles, velocity and acceleration
- Strain gauge readings (measure of load)

□ Full website at TU Dortmund, Germany

- http://www.rst.e-technik.tu-dortmund.de/cms/en/research/robotics/TUDOR_engl/index.html
- http://www.rst.e-technik.tu-dortmund.de/forschung/robot-toolbox/MERIt/MERIt_Documentation.pdf

