# Unit 8 Support Vector Machines

EL-GY 6143/CS-GY 6923: INTRODUCTION TO MACHINE LEARNING PROF. PEI LIU





#### Midterm

#### **■**Midterm

- March 29 (Monday), 11:00AM-1:30PM ET via Zoom
- Zoom link will be emailed later this week
- Exam questions via Google form
  - https://docs.google.com/forms/d/e/1FAIpQLSe8Ugqv7UhBv2B6zd7nYwhzX78cno7EnPjvfrdHaRK tBnBujQ/viewform?usp=sf link

#### Policy

- Close-book exam
- 2 piece of paper cheatsheet, can write on both sides.
- Attendance at exams is mandatory
- ■What will be included in the midterm?
  - All materials taught in Lecture 1 to Lecture 7
  - Not include Lecture 8: SVM
  - 6-7 questions, very similar to homework





#### Requirements for Midterm

- ☐ Please join the meeting on 15 minutes before 11AM, and use your full name when you join.
- □ Each student is required to open Zoom and turn on their video, making sure the camera captures your hands and your computer screen/keyboard. The whole exam will be recorded. To clearly capture the video of exam taking, it is recommended that:
  - A student can use an external webcam connected to the computer;
  - Or use another device (smartphone/ipad/laptop with power plugged in);
  - In both cases, adjust the position/orientation of the camera so that it clearly captures the keyboard, screen and both hands.
- □ During the exam, you should keep your video on all the time. If there is anything wrong with your Zoom connection, please recoonect ASAP. If you cannot reconnect, please email ASAP.





#### Midterm Submission

- □Students can read exam questions from computer screen, and write answers on BLANK papers using pen/pencil. If you have a printer, you can also print the questions and answer directly on the question papers.
  - Use a separate page for each question, and clear mark the questions on the top of the page.
- □ At 1:30PM, submit a single PDF to newclasses.nyu.edu, under the assignment named "Midterm Exam".
  - Use Adobe Scan APP/iPhone Notes APP/etc to scan your answers;
  - Don't upload several photos;
  - The DEADLINE for submission is in 10 minutes .If you never used your phone to scan a documents, give
    it a try before the exam;
  - Before you leave the exam, it's your responsibility to make sure all your answers are uploaded.





## Learning Objectives

- ☐ Interpret weights in linear classification of images
- ☐ Describe why linear classification for images does not work
- ☐ Define the margin in linear classification
- ☐ Describe the SVM classification problem.
- ☐ Write equations for solutions of constrained optimization using the Lagrangian.
- ☐ Describe a kernel SVM problem for non-linear classification
- ☐ Implement SVM classifiers in python
- ☐ Select SVM parameters from cross-validation





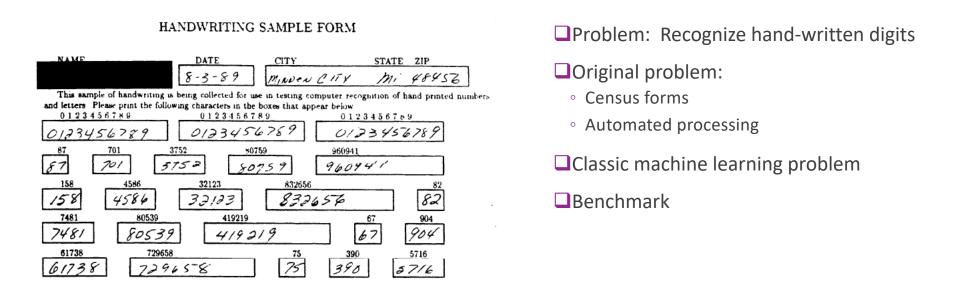
#### Outline

- Motivating example: Recognizing handwritten digits
  - Why logistic regression doesn't work well.
- ☐ Maximum margin classifiers
- ■Support vector machines
- ☐Kernel trick
- ☐ Constrained optimization





#### MNIST Digit Classification



From Patrick J. Grother, NIST Special Database, 1995



# A Widely-Used Benchmark

■We will look at SVM today

■ Not the best algorithm

☐But quite good

☐...and illustrates the main points

#### Classifiers [edit]

This is a table of some of the machine learning methods used on the database and their error rates, by type of classifier:

Type \$	Classifier +	Distortion +	Preprocessing +	Error rate (%) \$
Linear classifier	Pairwise linear classifier	None	Deskewing	7.6 <sup>[9]</sup>
K-Nearest Neighbors	K-NN with non-linear deformation (P2DHMDM)	None	Shiftable edges	0.52 <sup>[14]</sup>
Boosted Stumps	Product of stumps on Haar features	None	Haar features	0.87 <sup>[15]</sup>
Non-Linear Classifier	40 PCA + quadratic classifier	None	None	3.3[9]
Support vector machine	Virtual SVM, deg-9 poly, 2-pixel jittered	None	Deskewing	0.56 <sup>[16]</sup>
Neural network	2-layer 784-800-10	None	None	1.6 <sup>[17]</sup>
Neural network	2-layer 784-800-10	elastic distortions	None	0.7 <sup>[17]</sup>
Deep neural network	6-layer 784-2500-2000-1500-1000-500-10	elastic distortions	None	0.35 <sup>[18]</sup>
Convolutional neural network	Committee of 35 conv. net, 1-20-P-40-P-150-10	elastic distortions	Width normalizations	0.23[8]



## **Downloading MNIST**

```
import tensorflow as tf

(Xtr,ytr),(Xts,yts) = tf.keras.datasets.mnist.load_data()

print('Xtr shape: %s' % str(Xtr.shape))

print('Xts shape: %s' % str(Xts.shape))

ntr = Xtr.shape[0]

nts = Xts.shape[0]

nrow = Xtr.shape[1]

ncol = Xtr.shape[2]
```

Xtr shape: (60000, 28, 28) Xts shape: (10000, 28, 28)

- ■MNIST data is available in many sources
  - Note: It has been removed from sklearn
- ☐ Tensorflow version:
  - 60000 training samples
  - 10000 test samples
- ☐ Each sample is a 28 x 28 images
- □ Grayscale: Pixel values  $\in \{0,1,...,255\}$ 
  - ∘ 0 = Black and
  - 255 = White







#### Matrix and Vector Representation

- $\square$  For this demo, we reshape data from  $N \times 28 \times 28$  to  $N \times 784$
- ☐But, you can easily go back and forth
- □Also, scale the pixel values from -1 to 1







$$S = Mat(x) = \begin{bmatrix} s_{11} & \cdots & s_{1,28} \\ \vdots & \vdots & \vdots \\ s_{28,1} & \cdots & s_{28,28} \end{bmatrix}$$

$$x = \text{vec}(S) = \begin{bmatrix} x_1 & \cdots & x_{784} \end{bmatrix}$$

## Displaying Images in Python









4 random images in the dataset

A human can classify these easily

```
def plt_digit(x):
    nrow = 28
    ncol = 28
    xsq = x.reshape((nrow,ncol))
                                                 Key command
    plt.imshow(xsq, cmap='Greys_r') ←
    plt.xticks([])
    plt.yticks([])
# Convert data to a matrix
X = mnist.data
v = mnist.target
# Select random digits
                                                 Sample
nplt = 4
nsamp = X.shape[0]
                                                 permutation is
Iperm = np.random.permutation(nsamp)
                                                 necessary for this
# Plot the images using the subplot command
                                                 dataset, as the
for i in range(nplt):
                                                 original data is
    ind = Iperm[i]
    plt.subplot(1,nplt,i+1)
                                                 ordered by digits
    plt_digit(X[ind,:])
```



#### Try a Logistic Classifier

```
ntr1 = 5000
Xtr1 = Xtr[Iperm[:ntr1],:]
ytr1 = ytr[Iperm[:ntr1]]
```

- ☐ Train on 5000 samples
  - To reduce training time.
  - In practice want to train with ~40k
- Select correct solver (lbfgs)
  - Others can be very slow. Even this will take minutes

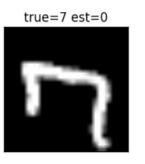


#### Performance

- □Accuracy = 89%. Very bad
- □Some of the errors seem like they should have been easy to spot
- ■What went wrong?

```
nts1 = 5000
Iperm_ts = np.random.permutation(nts)
Xts1 = Xts[Iperm_ts[:nts1],:]
yts1 = yts[Iperm_ts[:nts1]]
yhat = logreg.predict(Xts1)
acc = np.mean(yhat == yts1)
print('Accuaracy = {0:f}'.format(acc))
```

Accuaracy = 0.891000

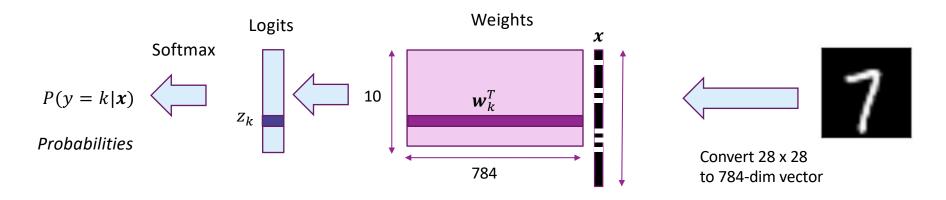








# Recap: Logistic Classifier

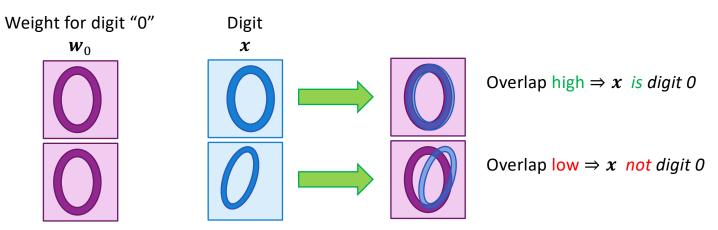


- □ Each logit  $z_k = \boldsymbol{w}_k^T \boldsymbol{x}$  = inner product with weight  $\boldsymbol{w}_k$  with digit  $\boldsymbol{x}$ , k = 0, ..., 9
- $\square \text{Will select } \hat{y} = \arg \max_{k} P(y = k | x) = \arg \max_{k} z_{k}$ 
  - $\circ$  Output  $z_k$  which is largest
- $\square$  When is  $z_k$  large?



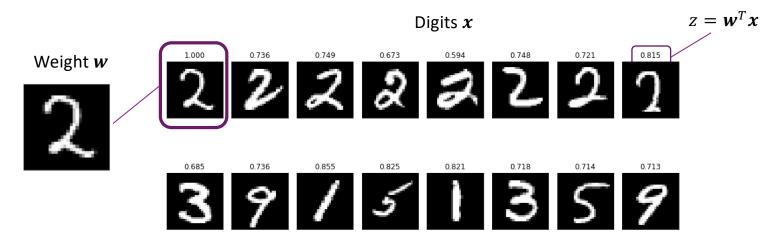
#### Interpreting the Logistic Classifier Weights

- $\square$  A logit  $z_k = w_k^T x$  is high when there is high overlap between  $w_k$  with digit x
  - Visualize each weight as an image
  - Suppose pixels are 0 or 1
  - $egin{aligned} & o & z_k = oldsymbol{w}_k^T oldsymbol{x} = \sum_i w_{ki} x_i = ext{number of pixels that overlap with } oldsymbol{w}_k ext{ and } oldsymbol{x} \end{aligned}$
- □Conclusion: Small variations in digits can cause low overlap



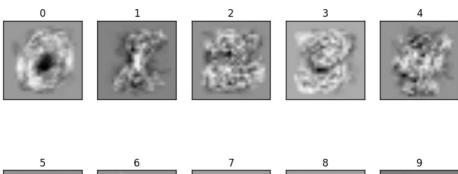
### **Example with Actual Digits**

- $\square$  Take weight w from a random digit "2"
- □Inner products  $z = \mathbf{w}^T \mathbf{x}$  are only slightly higher for other digits "2"
- $\Box$  Cannot tell which digit is correct from the inner product  $z = w^T x$



# Visualizing the Weights

- □Optimized weights of the classifier
- □Blurry versions of image to try to capture rotations, translations, ...













### Problems with Logistic Classifier

- ☐ Linear weighting cannot capture many deformities in image
  - Rotations
  - Translations
  - Variations in relative size of digit components
- ☐ Can be improved with preprocessing
  - E.g. deskewing, contrast normalization, many methods
- □ Is there a better classifier?





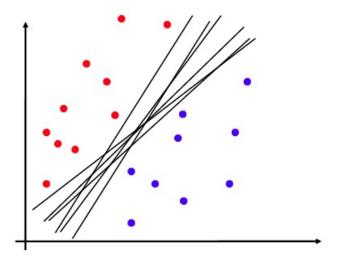
#### Outline

- ☐ Motivating example: Recognizing handwritten digits
  - Why logistic regression doesn't work well.
- Maximum margin classifiers
- ■Support vector machines
- ☐Kernel trick
- ☐ Constrained optimization



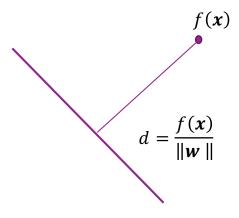
# Non-Uniqueness of Separating Plane

- ☐ Linearly separable data:
  - Can find a separating hyper-plane as a linear classifier.
- ☐ Separating hyper-plane is not unique
  - Fig. on right: Many separating planes
- ☐Which one is optimal?



## Hyperplane Basics

- □Linear function:  $f(x) = w^T x + b, x \in \mathbb{R}^d$
- $\square$  Hyperplane in d-dimensional: f(x) = 0
- □Parameters:
  - $\circ$  Weight w and bias b
  - Unique up to scaling:
  - $\circ$  (b, w) and  $(\alpha b, \alpha w)$  define the same plane.
  - $\circ$  For unique definition, we can require  $\| \boldsymbol{w} \|$ =1.
- □ Distance of any point **x** to the hyperplane:
  - d = f(x)/||w||, where  $f(x) = b + w^T x$ .
  - ∘ See ESL Sec. 4.5.
  - ESL: Hastie, Tibshirani, Friedman, "The Elements of Statistical Learning". 2<sup>nd</sup> Ed. Springer.



Hyperplane

$$f(\mathbf{x}) = \mathbf{w}^T \mathbf{x} + b = 0$$

# Linear Separability and Margin

- $\square$  Given training data  $(x_i, y_i)$ , i = 1, ..., N
  - Binary class label:  $y_i = \pm 1$
- $\square$ Suppose it is separable with parameters (w, b)
- □ There must exist a  $\gamma > 0$  s.t.:

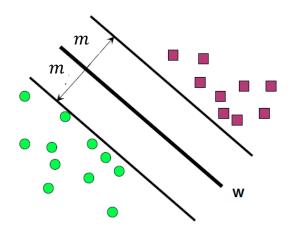
• 
$$b + w_1 x_{i1} + \cdots w_d x_{id} > \gamma$$
 when  $y_i = 1$ 

• 
$$b + w_1 x_{i1} + \cdots w_d x_{id} < -\gamma$$
 when  $y_i = -1$ 

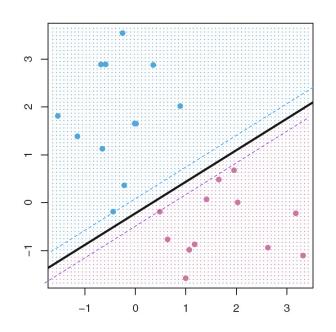
☐ Single equation form:

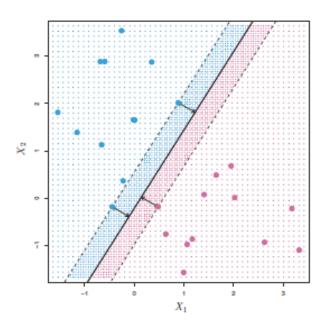
$$y_i(b + w_1x_{i1} + \cdots w_dx_{id}) > \gamma \text{ for all } i = 1, ..., N$$

- $\square$  Margin:  $m = \frac{\gamma}{\|w\|}$ : minimal distance of a sample to the plane
  - $^{\circ}$   $\gamma$  is the maximum value satisfying the above constraints



# Which separating plane is better?





From Fig. 9.2 and Fig. 9.3 in ISL.





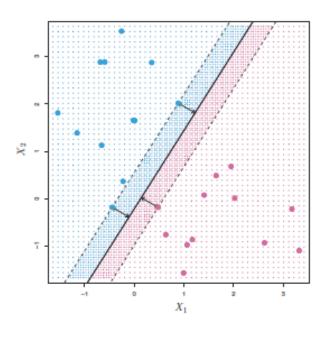
## Maximum Margin Classifier

- ☐ For the classifier to be more robust to noise, we want to maximize the margin!
- □ Define maximum margin classifier

$$\max_{w,b} \gamma$$
 Maximizes the margin 
$$\circ \text{ Such that } y_i(b+\textbf{w}^T\textbf{x}) \geq \gamma \text{ for all } i$$
 Ensures all points are correctly classified 
$$\sum_{j=1}^d w_j^2 \leq 1$$
 Scaling on weights

- □ Called a constrained optimization
  - Objective function and constraints
  - More on this later.
- See closed form solution in Sec. 4.5.2 in ESL. Note notation difference.

### Visualizing Maximum Margin Classifier



- ☐Fig. 9.3 of ISL
- ☐ Margin determined by closest points to the line
  - The maximal margin hyperplane represents the midline of the widest "slab" that we can insert between two classes
- ☐ In this figure, there are 3 points at the margin

ISL: James, Witten, Hastie, Tibshirani, An Introduction to Statistical Learning, Springer. 2013.

#### Problems with MM classifier

- ☐ Data is often not perfectly separable
  - Only want to correctly separate most points

- ■MM classifier is not robust
  - A single sample can radically change line

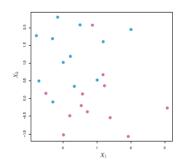


Fig. 9.4

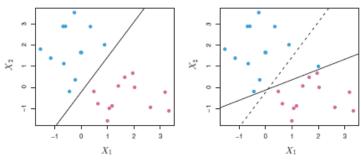


Fig. 9.5



#### **In-Class Exercise**

☐ Found in github site: svm\_inclass.ipynb

#### Problem 1. Margin

For the points below with binary labels:

- . Create a scatter plot of the points with different markers for the two classes
- . Find the weight and bias of the classifier that separates the two classes
- . Compute the distance to the classifier boundary for the points
- · Find the margin

```
X = np.array([[0.5,0.5], [1,0.5],[0.5,1.75], [0.75,2.75], [1.1,2.2], [2,1], [3,1.5]])

y = np.array([1,1,1,0,0,0,0])
```





#### Outline

- ☐ Motivating example: Recognizing handwritten digits
  - Why logistic regression doesn't work well.
- ☐ Maximum margin classifiers
- Support vector machines
  - ■Kernel trick
  - ☐ Constrained optimization

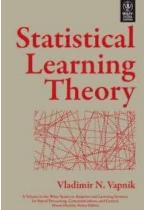




#### Support Vector Machine

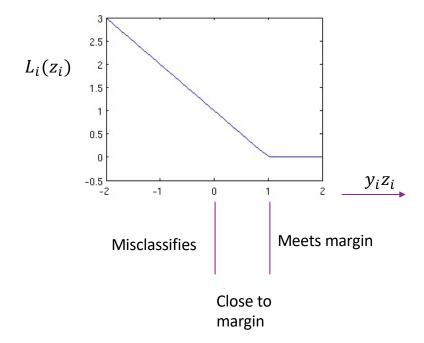
- ■Support Vector Machine (SVM)
  - Vladimir Vapnik, 1963
  - But became widely-used with kernel trick, 1993
  - More on this later
- ☐Got best results on character recognition
- ☐ Key idea: Allow "slack" in the classification
  - Support vector classifier (SVC): Directly use raw features.
     Good when the original feature space is roughly linearly separable
  - Support vector machine (SVM): Map the raw features to some other domain through a kernel function





### Hinge Loss

- $\Box$ Fix  $\gamma = 1$
- □Want ideally:  $y_i(\mathbf{w}^T\mathbf{x} + b) \ge 1$  for all samples i
  - Equivalently,  $y_i z_i \ge 1$ ,  $z_i = b + \mathbf{w}^T \mathbf{x}$
- ☐But perfect separation may not be possible
- □ Define hinge loss or soft margin:
  - $L_i(\mathbf{w}, b) = \max(0, 1 y_i z_i)$
- ☐Starts to increase as sample is misclassified:
  - $y_i z_i \ge 1 \implies$  Sample meets margin target,  $L_i(w) = 0$
  - $y_i z_i \in [0,1) \Rightarrow \text{Sample margin too small, small loss}$
  - ∘  $y_i z_i \le 0$  ⇒ Sample misclassified, large loss



# **SVM Optimization**

- $\square$  Given data  $(x_i, y_i)$

$$J(\boldsymbol{w}, b) = C \sum_{i=1}^{N} \max(0, 1 - y_i(\boldsymbol{w}^T \boldsymbol{x}_i + b)) + \frac{1}{2} \|\boldsymbol{w}\|^2$$

$$\text{margin} \quad \text{Hinge loss term} \quad \text{margin} = 1/\|\boldsymbol{w}\|$$

C controls final margin

Hinge loss term
Attempts to reduce
Misclassifications

- $\square$ Constant C > 0 will be discussed below
- Note: ISL book uses different naming conventions.
  - We have followed convention in sklearn

### Alternate Form of SVM Optimization

☐ Equivalent optimization:

$$\min J_1(\boldsymbol{w}, b, \boldsymbol{\epsilon}), \qquad J_1(\boldsymbol{w}, b, \boldsymbol{\epsilon}) = C \sum_{i=1}^N \epsilon_i + \frac{1}{2} \|\boldsymbol{w}\|^2$$

■ Subject to constraints:

$$y_i(\mathbf{w}^T \mathbf{x}_i + b) \ge 1 - \epsilon_i \text{ for all } i = 1, ..., N$$
  
 $\epsilon_i \ge 0 \text{ for all } i = 1, ..., N$ 

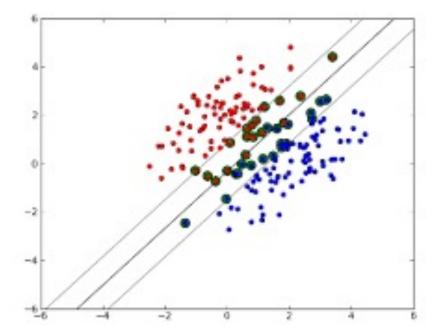
- $\epsilon_i$  = amount sample i misses margin target
- $\square$  Sometimes write as  $J_1(\boldsymbol{w}, b, \epsilon) = C \|\epsilon\|_1 + \frac{1}{2} \|\boldsymbol{w}\|^2$ 
  - $\| \epsilon \|_1 = \sum_{i=1}^N \epsilon_i \; ext{ called the "one-norm"}$
  - Generally one-norm would have absolute sign over  $\epsilon_i$ .
  - But in this case, when the constraint is met,  $\epsilon_i$ >=0.



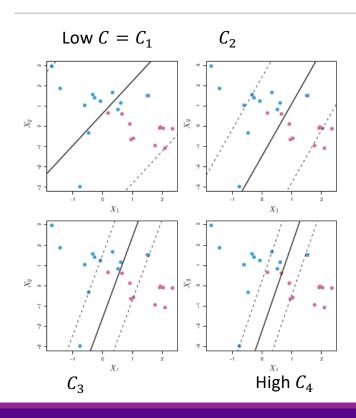


### **Support Vectors**

- □Support vectors: Samples that either:
  - Are exactly on margin:  $y_i(\mathbf{w}^T \mathbf{x}_i + b) = 1$
  - $\circ$  Or, on wrong side of margin:  $y_i(\mathbf{w}^T\mathbf{x}_i+b) \leq 1$
- ☐ Changing samples that are not SVs
  - Does not change solution
  - Provides robustness



# Illustrating Effect of C



#### ☐Fig. 9.7 of ISL

- Note: C has opposite meaning in ISL than python
- Here, we use python meaning

#### $\square$ Low C:

- Leads to large margin
- But allow many violations of margin.
- Many more SVs
- Reduces variance by using more samples

#### ☐ Large C:

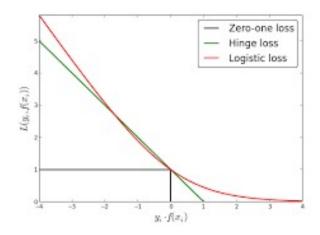
- Leads to small margin
- Reduce number of violations, and fewer SVs.
- Highly fit to data. Low bias, higher variance
- More chance to overfit



# Relation to Logistic Regression

□ Logistic regression also minimizes a loss function:

$$J(\mathbf{w}, b) = \sum_{i=1}^{N} L_i(\mathbf{w}, b), \qquad L_i(\mathbf{w}, b) = \ln P(y_i | \mathbf{x}_i) = -\ln(1 + e^{-y_i z_i})$$





#### **In-Class Exercise**

#### **Problem 2. Minimizing the Hinge Loss**

For the data below, first create a scatter plot of the points with different markers for the two classes. You should see that the data is not linearly separable.

Then, consider a set of classifiers:

```
yhat = sign(z), z = w.dot(x)+b
```

Use the the w below, plot the hinge loss as a function of the bias b where the hinge loss is:

```
J = sum( maximum(0, 1-ypm*z) )
```

Here ypm=2\*y-1 so that it is a value +1 or -1. Find the b that minimizes the hinge loss and plot the boundary of the classifier.

```
X = np.array([[0.5,0.5], [1,0.5],[0.5,1.75], [2,2], [0.75,0.75], [0.75,2.75], [1.1,2.2], [2,1], [3,1.5]])
y = np.array([1,1,1,1,0,0,0,0,0])

w = np.array([1.5, 1])
w = w / np.linalg.norm(w)
```



## Outline

- ☐ Motivating example: Recognizing handwritten digits
  - Why logistic regression doesn't work well.
- ☐ Maximum margin classifiers
- ■Support vector machines
- Kernel trick
- ☐ Constrained optimization



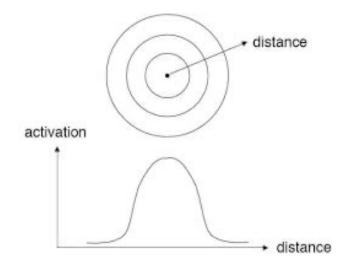
## The Kernel Function

#### ■Kernel function:

- Function  $K(x_i, x)$
- Key function for SVMs and kernel classifiers
- $^{\circ}$  Measures "similarity" between new sample  $oldsymbol{x}$  and training sample  $oldsymbol{x}_i$

#### ☐ Typical property

- $\cdot x_i, x \text{ close} \Rightarrow K(x_i, x) \text{ maximum value}$
- $x_i, x \text{ far} \Rightarrow K(x_i, x) \approx 0$



## Common Kernels

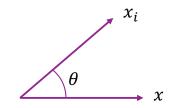
#### ☐Linear SVM:

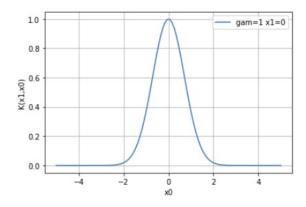
- $\circ K(x_i, x) = x_i^T x = ||x_i|| ||x|| \cos \theta$
- Maximum when angle between vectors is small
- □ Radial basis function:

$$K(x_i, x) = \exp[-\gamma ||x - x_i||^2]$$

 $\circ~1/\gamma$  indicates width of kernel

• Typically d=2

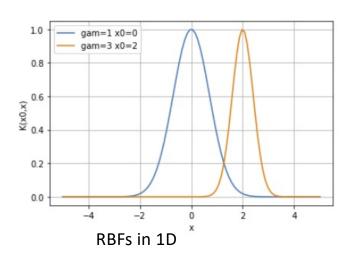


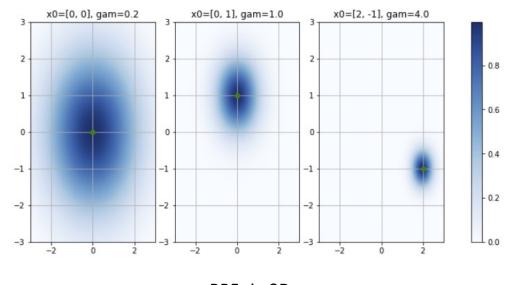


# **RBF Kernel Examples**

□RBF kernel:  $K(x_0, x) = \exp[-\gamma ||x - x_0||^2]$ 

- $\circ$  Peak value of 1 at  $x=x_0$
- Width  $\propto \frac{1}{\gamma}$









## Kernel Classifier

#### **□**Given:

- Training data  $(x_i, y_i)$  with binary labels  $y_i = \pm 1$
- Kernel  $K(x_i, x)$

#### $\blacksquare$ To classify a new point x:

- Decision function:  $z = \sum_{i=1}^{n} y_i K(x_i, x)$
- Classify:  $\hat{y} = sign(z)$

#### □Idea:

- $\circ z$  is large positive when x is close to samples  $x_i$  with  $y_i = 1$
- $\circ \ z$  is large negative when x is close to samples  $x_i$  with  $y_i = -1$
- ☐ Kernel classifiers are a subject on their own
  - We just mention them here to explain connection to SVMs

# Example in 1D

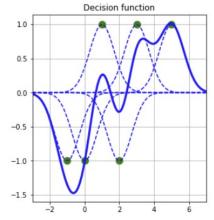
- $\blacksquare$  Example data with 6 points  $(x_i, y_i)$ 
  - RBF kernel:  $K(x_i, x) = e^{-\gamma(x_i x)^2}$ ,  $\gamma = 1$
- □ Decision function:

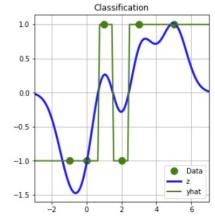
$$\circ z = \sum_{i=1}^n y_i K(x_i, x)$$

- Sum of bell curves
- Positive when near positive samples
- Negative when near negative samples
- **Classification**:

$$\circ \hat{y} = sign(z)$$


	i	1	2	3	4	5	6
2	$x_i$	-1	0	1	2	3	5
3	Vi	-1	-1	1	-1	1	1





## **Effect of Gamma**

- ■Same data as before
- $\square RBF kernel: K(x_i, x) = e^{-\gamma(x_i x)^2}$
- $\square$  As  $\gamma$  increases:
  - $\circ$  Decision function  $z \approx y_i$  when  $x = x_i$
  - Classifier fits training data better
  - Classification region more complex
- $\square$  As a classifier, higher  $\gamma$  results in:
  - Lower bias error
  - But, higher variance error



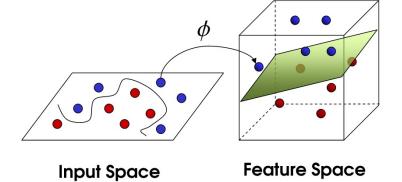
## **SVMs with Non-Linear Transformations**

#### ■ Non-linear transformation:

- Replace x with  $\phi(x)$
- Enables more rich, non-linear classifiers
- Examples: polynomial classification

$$\phi(x) = [1, x, x^2, \dots, x^{d-1}]$$

☐ Tries to find separation in a feature space



- □Kernel trick in SVMs:
  - Makes applying non-linear transformations easy

### SVM with the Transformation

- $\square$  Consider SVM model with x replaced by  $\phi(x)$
- ☐ Minimize SVM cost function as before (i.e. Hinge loss + inverse margin)
- ☐ Theorem: The optimal weight is of the form:

$$\mathbf{w} = \sum_{i=1}^{N} \alpha_i y_i \phi(\mathbf{x}_i)$$

- $\alpha_i \geq 0$  for all i
- $\alpha_i > 0$  if and only if sample i is a support vector
- Will show this fact later using results in constrained optimization
- $\square$  Consequence: The linear discriminant on any other sample x is:

$$z = b + \mathbf{w}^T \phi(\mathbf{x}) = b + \sum_{i=1}^N \alpha_i y_i \boxed{\phi(\mathbf{x}_i)^T \phi(\mathbf{x})} - K(\mathbf{x}_i, \mathbf{x}) = \text{"kernel"}$$



## Kernel Form of the SVM Classifier

□SVM classifier can be written with the kernel  $K(x_i, x)$  and values  $\alpha_i \ge 0$ :

$$z = b + \sum_{i=1}^{N} \alpha_i y_i K(x_i, x),$$

$$\hat{y} = \text{sign}(z) = \begin{cases} 1 & \text{if } z > 0 \\ -1 & \text{if } z < 0 \end{cases}$$
Classification decision

- □ Key point: SVM classifier is approximately Kernel classifier
- ■But there are two differences:
  - Weights  $\alpha_i \geq 0$  on the samples (the weights are only non-zero on the SVs)
  - A bias term b (can be positive or negative)

## "Kernel Trick" and Dual Parameterization

□ Kernel form of SVM classifier (previous slide):

$$z = b + \sum_{i=1}^{N} \alpha_i y_i K(\mathbf{x}_i, \mathbf{x}),$$
  
$$\hat{y} = \operatorname{sign}(z)$$

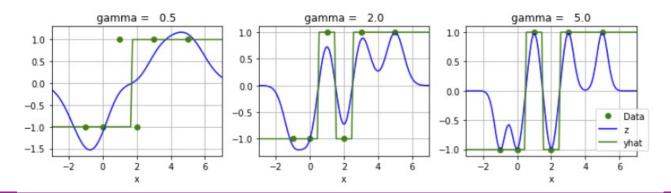
- □ Dual parameters:  $\alpha_i \ge 0, i = 1, ..., N$ 
  - Called the dual parameters due to constrained optimization see next section
- ☐Kernel trick:
  - $^{\circ}$  Directly solve the parameters lpha instead of the weights w
  - $\circ$  Can show that the optimization only needs the kernel  $K(x_i, x)$
  - $\circ$  Does not need to explicitly use  $\phi(x)$



# SVM Example in 1D


- ☐ Same data as in the Kernel classifier example
- $\square$  Fit SVM with RBF with different  $\gamma$
- $\square$  Similar trends as kernel classifier: As  $\gamma$  increases
  - z "fits" data  $(x_i, y_i)$  closer
  - Leads to more complex decision regions.
  - Enables nonlinear decision regions

i	1	2	3	4	5	6
$x_i$	-1	0	1	2	3	5
$y_i$	-1	-1	1	-1	1	1





# Example in 2D



#### ■Example:

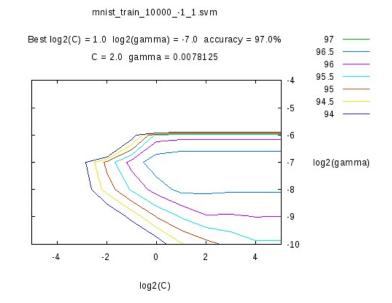
- 10 data points with binary labels
- $^{\circ}$  Fit SVM with C=1 and RBF
- $\gamma = 0.3, 3 \text{ and } 10$

#### ■Plot:

- $\circ$  z= linear discriminant
- $\hat{y} = sign(z) = classification decision$
- $\square$  Observe: As  $\gamma$  increases
  - Fits training data better
  - More complex decision region

### Parameter Selection

- ☐ For SVMs with RBFs we need to select:
  - Parameter C > 0 in the loss function
  - $\circ$  Kernel width  $\gamma > 0$
- $\square$  Higher C or  $\gamma$ 
  - Fewer SVs
  - Classifiers averages over smaller set
  - Lower bias, but higher variance
- ☐ Typically select via cross-validation
  - Try out different  $(C, \gamma)$  pairs
  - Find which one provides highest accuracy on test set
- ☐ Python can automatically do grid search



http://peekaboo-vision.blogspot.com/2010/09/mnist-for-ever.html





## Multi-Class SVMs

- $\square$ Suppose there are K classes
- One-vs-one:
  - Train  $\binom{K}{2}$  SVMs for each pair of classes
  - Test sample assigned to class that wins "majority of votes"
  - Best results but very slow
- One-vs-rest:
  - $\circ$  Train K SVMs: train each class k against all other classes
  - $\circ$  Pick class with highest  $z_k$
- ■Sklearn has both options





## **MNIST** Results

- ☐ Run classifier
- ■Very slow
  - Several minutes for 40,000 samples
  - Slow in training and test
  - Major drawback of SVM
- $\square$ Accuracy  $\approx 0.984$ 
  - Much better than logistic regression
- □Can get better with:
  - pre-processing
  - More training data
  - Optimal parameter selection

```
from sklearn import svm
# Create a classifier: a support vector classifier
svc = svm.SVC(probability=False, kernel="rbf", C=2.8, gamma=.0073,verbose=10)
```

```
[LibSVM]

SVC(C=2.8, cache_size=200, class_weight=None, coef0=0.0,
  decision_function_shape=None, degree=3, gamma=0.0073, kernel='rbf',
  max_iter=-1, probability=False, random_state=None, shrinking=True,
  tol=0.001, verbose=10)
```

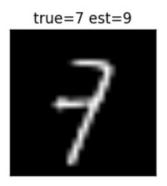
```
yhat1 = svc.predict(Xts)
acc = np.mean(yhat1 == yts)
print('Accuaracy = {0:f}'.format(acc))
```

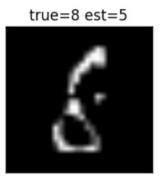
svc.fit(Xtr,ytr)

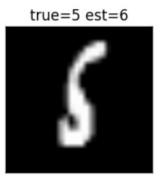
Accuaracy = 0.984000

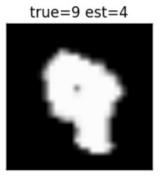
## **MNIST Errors**

■Some of the error are hard even for a human









## Outline

- ☐ Motivating example: Recognizing handwritten digits
  - Why logistic regression doesn't work well.
- ☐ Maximum margin classifiers
- ■Support vector machines
- ☐Kernel trick
- Constrained optimization



# **Constrained Optimization**

- ☐ In many problems, variables are constrained
- □ Constrained optimization formulation:
  - Objective: Minimize f(w)
  - Constraints:  $g_1(\mathbf{w}) \leq 0, ..., g_M(\mathbf{w}) \leq 0$
- ■Examples:
  - Minimize the mpg of a car subject to a cost or meeting some performance
  - In ML: weight vector may have constraints from physical knowledge
- $\square$  Often write constraints in vector form: Write  $g(w) \leq 0$

$$g(\mathbf{w}) = [g_1(\mathbf{w}), ..., g_m(\mathbf{w})]^T$$



# Lagrangian

- □ Constrained optimization: Min f(w) s.t.  $g(w) \le 0$
- $\square$  Consider first a single constraint: g(w) is a scalar
- □ Define Lagrangian:  $L(\mathbf{w}, \lambda) = f(\mathbf{w}) + \lambda g(\mathbf{w})$ 
  - w is called the primal variable
  - $\circ$   $\lambda$  is called the dual variable
- $\square$  Dual minimization: Given a dual parameter  $\lambda$ , minimize

$$\widehat{\boldsymbol{w}}(\lambda) = \arg\min_{\boldsymbol{w}} L(\boldsymbol{w}, \lambda), \qquad L^*(\lambda) = \min_{\boldsymbol{w}} L(\boldsymbol{w}, \lambda)$$

- Minimizes a weighted combination of objective and constraint.
- Higher  $\lambda \Rightarrow$  Weight constraint more (try to make  $g(\mathbf{w})$  smaller)
- Lower  $\lambda \Rightarrow$  Weight objective more (try to make f(w) smaller)

## **KKT Conditions**

- $\square$  Given objective f(w) and constraint g(w)
- $\square$ KKT Conditions:  $\widehat{w}$ ,  $\widehat{\lambda}$  satisfy:
  - $\widehat{\boldsymbol{w}}$  minimizes the Lagrangian:  $\widehat{\boldsymbol{w}} = \arg\min_{\boldsymbol{w}} L(\boldsymbol{w}, \widehat{\lambda})$
  - Either
    - $g(\widehat{\mathbf{w}}) = 0$  and  $\widehat{\lambda} \ge 0$  [active constraint]
    - $g(\widehat{\pmb{w}}) < 0$  and  $\widehat{\lambda} = 0$  [inactive constraint]
- ☐ Theorem: Under some technical conditions,
  - $\circ$  if  $\hat{w}$ ,  $\hat{\lambda}$  are local mimima of the constrained optimization, they must satisfy KKT conditions



# General Procedure for Single Constraint

#### ■Suppose:

- $\mathbf{w} = (w_1, ..., w_d)^T$ : d unknown primal variables
- $g(\mathbf{w}) \leq 0$ : scalar constraint

#### □Case 1: Assume constraint is active:

- Solve  $\mathbf{w}$  and  $\lambda$ :  $\partial L(\mathbf{w}, \lambda)/\partial w_i = 0$  and  $g(\mathbf{w}) = 0$  (resulting from setting  $\partial L(\mathbf{w}, \lambda)/\partial \lambda = 0$ )
- $\circ d + 1$  unknowns and d + 1 equations
- Verify that  $\lambda \geq 0$

#### □ Case 2: Assume constraint is inactive

- Solve primal objective  $\partial f(\mathbf{w})/\partial w_i = 0$  ignoring constraint
- $\circ \ d$  unknowns and d equations
- Verify that constraint is satisfied:  $g(\mathbf{w}) \leq 0$





## KKT Conditions Illustrated

☐ Example 1: Constraint is "active"

$$\min_{w} w^2 \quad s. t. \ w + 1 \le 0$$

■ Example 2: Constraint is "inactive"

$$\min_{w} w^2 \quad s. t. \ w - 1 \le 0$$

☐ Examples worked on board with illustration

# Multiple Constraints

- □ Now consider constraint:  $g(\mathbf{w}) = [g_1(\mathbf{w}), ..., g_M(\mathbf{w})]^T \le 0$ .
- ☐ Lagrangian is:

$$L(\mathbf{w}, \lambda) = f(\mathbf{w}) + \lambda^T g(\mathbf{w}) = f(\mathbf{w}) + \sum_{m=1}^{M} \lambda_m g_m(\mathbf{w})$$

- Weighted sum of all M constraints
- $\circ$   $\lambda$  is called the dual vector
- □KKT conditions extend to:
  - $\widehat{\boldsymbol{w}}$  minimizes the Lagrangian:  $\widehat{\boldsymbol{w}} = \arg\min_{\boldsymbol{w}} L(\boldsymbol{w}, \widehat{\lambda})$
  - $\circ$  For each  $m=1,\ldots,M$ 
    - $g_m(\widehat{\pmb{w}}) = 0$  and  $\hat{\lambda}_m \geq 0$  [active constraint]
    - $g_m(\widehat{\pmb{w}}) < 0$  and  $\hat{\lambda}_m = 0$  [inactive constraint]



# Multiple Constraints

- $\square$  If there are M constraints, there could be  $2^M$  cases to discuss
- ☐ In practice, the number of cases are much smaller
- ☐ For more information on KKT conditions, check the following lecture on youtube
  - <u>UAMathCamp</u> Lecture 40(A): Kuhn-Tucker Conditions: Conceptual and geometric insight
  - https://www.youtube.com/watch?v=HIm3Z0L90Co



# **SVM Constrained Optimization**

□ Recall: SVM constrained optimization

$$\min J_1(\boldsymbol{w}, b, \boldsymbol{\epsilon}), \qquad J_1(\boldsymbol{w}, b, \boldsymbol{\epsilon}) = C \sum_{i=1}^N \epsilon_i + \frac{1}{2} \|\boldsymbol{w}\|^2$$

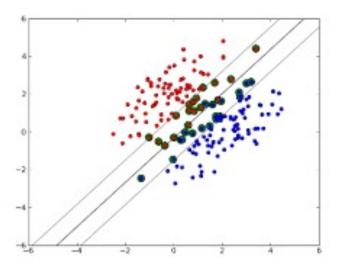
- Constraints:  $y_i(\mathbf{w}^T \mathbf{x}_i + b) \ge 1 \epsilon_i$  and  $\epsilon_i \ge 0$  for all i = 1, ..., N
- ☐ After applying KKT conditions and some algebra [beyond this class], solution is
  - $\circ$  Optimal weight vector:  $m{w} = \sum_{i=1}^N lpha_i y_i m{x}_i$  linear combination of instances
  - $\circ$  Dual parameters  $\alpha_i$  minimize

$$\sum_{i=1}^{N} \alpha_i - \frac{1}{2} \sum_{i=1}^{N} \sum_{j=1}^{N} \alpha_i \alpha_j y_i y_j \mathbf{x}_i^T \mathbf{x}_j \quad \text{s.t. } 0 \le \alpha_i \le C$$

• Details can be found on textbook ESLII, section 12.2.1 Computing the Support Vector Classifier

# **Support Vectors**

- $\square$  Classifier weight is:  $\mathbf{w} = \sum_{i=1}^{N} \alpha_i y_i \mathbf{x}_i$
- $\square$  Can show that  $\alpha_i > 0$  only when  $x_i$  is a support vector
  - On boundary or violating constraint
  - $\circ$  Otherwise  $\alpha_i=0$



## What you should know

- □ Interpret weights in linear classification of images (logistic regression): Match filters
- ☐ Understand the margin in linear classification and maximum margin classifier
- □SVM classifier: Allow violation of margin by introducing slack variables (More robust than linear classifier)
- □ Solve constrained optimization using the Lagrangian.
  - Understand KKT conditions for a single constraint
- ■Extend to nonlinear classifier by feature transformation: SVM with nonlinear kernels
- ☐ Select SVM parameters from cross-validation



